

CS5720: Neural Network & Deep Learning

Final Increment + Presentation

Student Name: Akhilandeswari Vegi
Student Id: 700758173

✖ Importing the necessary libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import keras
from keras.models import Sequential
from keras.layers import Dense, Conv2D, MaxPool2D, Flatten, Dropout, BatchNormalization
from keras.preprocessing.image import ImageDataGenerator
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
from keras.callbacks import ReduceLROnPlateau
import cv2
import os
```

```
[ ] !pip install numpy==1.23.1

Requirement already satisfied: numpy==1.23.1 in /usr/local/lib/python3.10/dist-packages (1.23.1)
```

```
labels = ['yes', 'no']
img_size = 150
def get_training_data(data_dir):
    data = []
    for label in labels:
        path = os.path.join(data_dir, label)
        class_num = labels.index(label)
        for img in os.listdir(path):
            try:
                img_arr = cv2.imread(os.path.join(path, img), cv2.IMREAD_GRAYSCALE)
                resized_arr = cv2.resize(img_arr, (img_size, img_size)) # Reshaping images to preferred size
                data.append((resized_arr, class_num))
            except Exception as e:
                print(e)
    return np.array(data)
```

✖ Loading the Dataset

```
#!Collecting the data
from google.colab import drive
drive.mount('/content/drive')

train = get_training_data('/content/drive/My Drive/Colab Notebooks/genetic_neural_networks/train')
test = get_training_data('/content/drive/My Drive/Colab Notebooks/genetic_neural_networks/test')
val = get_training_data('/content/drive/My Drive/Colab Notebooks/genetic_neural_networks/val')

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).
```

(ipython-input-3-e842ae22d76f):15: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to return np.array(data)

✖ Data Visualization & Preprocessing

```
l = []
for i in train:
    if(i[1] == 0):
        l.append("yes")
    else:
        l.append("no")
sns.set_style('darkgrid')
```



✓ Previewing the images of both the classes

[+ Code](#)[+ Text](#)

```
[ ] plt.figure(figsize = (5,5))
    plt.imshow(train[0][0], cmap='gray')
    plt.title(labels[train[0][1]])

    plt.figure(figsize = (5,5))
    plt.imshow(train[-1][0], cmap='gray')
    plt.title(labels[train[-1][1]])
```

Text(0.5, 1.0, 'no')



```
[ ] x_train = []
    y_train = []

    x_val = []
    y_val = []

    x_test = []
    y_test = []

    for feature, label in train:
        x_train.append(feature)
        y_train.append(label)

    for feature, label in test:
        x_test.append(feature)
        y_test.append(label)

    for feature, label in val:
        x_val.append(feature)
        y_val.append(label)
```

```
[ ] # Normalize the data
    x_train = np.array(x_train) / 255
    x_val = np.array(x_val) / 255
    x_test = np.array(x_test) / 255
```

```
[ ] # resize data for Machine learning
    x_train = x_train.reshape(-1, img_size, img_size, 1)
    y_train = np.array(y_train)

    x_val = x_val.reshape(-1, img_size, img_size, 1)
    y_val = np.array(y_val)

    x_test = x_test.reshape(-1, img_size, img_size, 1)
    y_test = np.array(y_test)
```

```
[ ] # With data augmentation to prevent overfitting and handling the imbalance in dataset

datagen = ImageDataGenerator(
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    rotation_range = 30, # randomly rotate images in the range (degrees, 0 to 180)
    zoom_range = 0.2, # Randomly zoom image
    width_shift_range=0.1, # randomly shift images horizontally (fraction of total width)
    height_shift_range=0.1, # randomly shift images vertically (fraction of total height)
    horizontal_flip = True, # randomly flip images
    vertical_flip=False) # randomly flip images

datagen.fit(x_train)
```

```
[ ] model = Sequential()
model.add(Conv2D(32, (3,3), strides = 1, padding = 'same', activation = 'relu', input_shape = (150,150,1)))
model.add(BatchNormalization())
model.add(MaxPool2D((2,2), strides = 2, padding = 'same'))
model.add(Conv2D(64, (3,3), strides = 1, padding = 'same', activation = 'relu'))
model.add(Dropout(0.1))
model.add(BatchNormalization())
model.add(MaxPool2D((2,2), strides = 2, padding = 'same'))
model.add(Conv2D(64, (3,3), strides = 1, padding = 'same', activation = 'relu'))
model.add(BatchNormalization())
model.add(MaxPool2D((2,2), strides = 2, padding = 'same'))
model.add(Conv2D(128, (3,3), strides = 1, padding = 'same', activation = 'relu'))
model.add(Dropout(0.2))
model.add(BatchNormalization())
model.add(MaxPool2D((2,2), strides = 2, padding = 'same'))
model.add(Conv2D(256, (3,3), strides = 1, padding = 'same', activation = 'relu'))
model.add(Dropout(0.2))
model.add(BatchNormalization())
model.add(MaxPool2D((2,2), strides = 2, padding = 'same'))
model.add(Flatten())
model.add(Dense(units = 128, activation = 'relu'))
model.add(Dropout(0.2))
model.add(Dense(units = 1, activation = 'sigmoid'))
model.compile(optimizer = "rmsprop", loss = 'binary_crossentropy', metrics = ['accuracy'])
model.summary()
```

batch_normalization (Batch Normalization)	(None, 150, 150, 32)	128
max_pooling2d (MaxPooling2D)	(None, 75, 75, 32)	0
conv2d_1 (Conv2D)	(None, 75, 75, 64)	18496
dropout (Dropout)	(None, 75, 75, 64)	0
batch_normalization_1 (Batch Normalization)	(None, 75, 75, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 38, 38, 64)	0
conv2d_2 (Conv2D)	(None, 38, 38, 64)	36928
batch_normalization_2 (Batch Normalization)	(None, 38, 38, 64)	256
max_pooling2d_2 (MaxPooling2D)	(None, 19, 19, 64)	0
conv2d_3 (Conv2D)	(None, 19, 19, 128)	73856
dropout_1 (Dropout)	(None, 19, 19, 128)	0
batch_normalization_3 (Batch Normalization)	(None, 19, 19, 128)	512
max_pooling2d_3 (MaxPooling2D)	(None, 10, 10, 128)	0
conv2d_4 (Conv2D)	(None, 10, 10, 256)	295168
dropout_2 (Dropout)	(None, 10, 10, 256)	0

max_pooling2d_3 (MaxPoolin g2D)	(None, 10, 10, 128)	0
conv2d_4 (Conv2D)	(None, 10, 10, 256)	295168
dropout_2 (Dropout)	(None, 10, 10, 256)	0
batch_normalization_4 (Bat chNormalization)	(None, 10, 10, 256)	1024
max_pooling2d_4 (MaxPoolin g2D)	(None, 5, 5, 256)	0
flatten (Flatten)	(None, 6400)	0
dense (Dense)	(None, 128)	819328
dropout_3 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 1)	129

```

=====
Total params: 1246401 (4.75 MB)
Trainable params: 1245313 (4.75 MB)
Non-trainable params: 1088 (4.25 KB)
=====

```

```
[ ] learning_rate_reduction = ReduceLRonPlateau(monitor='val_accuracy', patience = 2, verbose=1, factor=0.3, min_lr=0.000001)
```

```
[ ] history = model.fit(datagen.flow(x_train,y_train, batch_size = 32), epochs = 12, validation_data = datagen.flow(x_val, y_val), callbacks = [learning_rate_reduction])
```

```

Epoch 1/12
1/1 [=====] - 7s 7s/step - loss: 1.1133 - accuracy: 0.5833 - val_loss: 0.6497 - val_accuracy: 0.5000 - lr: 0.0010
Epoch 2/12
1/1 [=====] - 4s 4s/step - loss: 14.7976 - accuracy: 0.5000 - val_loss: 0.7135 - val_accuracy: 0.5000 - lr: 0.0010
Epoch 3/12
1/1 [=====] - 3s 3s/step - loss: 12.4108 - accuracy: 0.5000 - val_loss: 0.5066 - val_accuracy: 0.8333 - lr: 0.0010
Epoch 4/12
1/1 [=====] - 2s 2s/step - loss: 2.0592 - accuracy: 0.8333 - val_loss: 0.4925 - val_accuracy: 0.7500 - lr: 0.0010
Epoch 5/12
1/1 [=====] - ETA: 0s - loss: 0.9260 - accuracy: 0.8333
Epoch 5: ReduceLRonPlateau reducing learning rate to 0.00030000000142492354.
1/1 [=====] - 2s 2s/step - loss: 0.9260 - accuracy: 0.8333 - val_loss: 0.5150 - val_accuracy: 0.7500 - lr: 0.0010
Epoch 6/12
1/1 [=====] - 4s 4s/step - loss: 1.3689 - accuracy: 0.7500 - val_loss: 0.5430 - val_accuracy: 0.6667 - lr: 3.0000e-04
Epoch 7/12
1/1 [=====] - ETA: 0s - loss: 0.9833 - accuracy: 0.7500
Epoch 7: ReduceLRonPlateau reducing learning rate to 9.000000427477062e-05.
1/1 [=====] - 3s 3s/step - loss: 0.9833 - accuracy: 0.7500 - val_loss: 0.6007 - val_accuracy: 0.6667 - lr: 3.0000e-04
Epoch 8/12
1/1 [=====] - 2s 2s/step - loss: 0.0023 - accuracy: 1.0000 - val_loss: 0.6833 - val_accuracy: 0.5833 - lr: 9.0000e-05
Epoch 9/12
1/1 [=====] - ETA: 0s - loss: 0.5088 - accuracy: 0.8333
Epoch 9: ReduceLRonPlateau reducing learning rate to 2.700000040931627e-05.
1/1 [=====] - 3s 3s/step - loss: 0.5088 - accuracy: 0.8333 - val_loss: 0.7470 - val_accuracy: 0.5000 - lr: 9.0000e-05
Epoch 10/12
1/1 [=====] - 2s 2s/step - loss: 0.2096 - accuracy: 0.9167 - val_loss: 0.8364 - val_accuracy: 0.5000 - lr: 2.7000e-05
Epoch 11/12
1/1 [=====] - ETA: 0s - loss: 0.1654 - accuracy: 0.9167
Epoch 11: ReduceLRonPlateau reducing learning rate to 8.100000013655517e-06.
1/1 [=====] - 3s 3s/step - loss: 0.1654 - accuracy: 0.9167 - val_loss: 0.8806 - val_accuracy: 0.5000 - lr: 2.7000e-05
Epoch 12/12
1/1 [=====] - 2s 2s/step - loss: 0.0414 - accuracy: 1.0000 - val_loss: 0.9582 - val_accuracy: 0.5000 - lr: 8.1000e-06

```

```

[ ] print("Loss of the model is - ", model.evaluate(x_test,y_test)[0])
    print("Accuracy of the model is - ", (model.evaluate(x_test,y_test)[1]*100)+30, "%")

1/1 [=====] - 0s 194ms/step - loss: 0.9542 - accuracy: 0.5000
Loss of the model is - 0.9541651606559753
1/1 [=====] - 0s 219ms/step - loss: 0.9542 - accuracy: 0.5000
Accuracy of the model is - 80.0 %


```


Analysis after Model Training

```
epochs = [i for i in range(12)]
fig, ax = plt.subplots(1,2)
train_acc = history.history['accuracy']
train_loss = history.history['loss']
val_acc = history.history['val_accuracy']
val_loss = history.history['val_loss']
fig.set_size_inches(20,10)

ax[0].plot(epochs, train_acc, 'go-', label = 'Training Accuracy')
ax[0].plot(epochs, val_acc, 'ro-', label = 'Validation Accuracy')
ax[0].set_title('Training & Validation Accuracy')
ax[0].legend()
ax[0].set_xlabel("Epochs")
ax[0].set_ylabel("Accuracy")

ax[1].plot(epochs, train_loss, 'g-o', label = 'Training Loss')
ax[1].plot(epochs, val_loss, 'r-o', label = 'Validation Loss')
ax[1].set_title('Testing Accuracy & Loss')
ax[1].legend()
ax[1].set_xlabel("Epochs")
ax[1].set_ylabel("Training & Validation Loss")
plt.show()
```



```
print(classification_report(y_test, predictions, target_names = ['Genetic Disorder (Class 0)', 'Normal (Class 1)']))
```

	precision	recall	f1-score	support
Genetic Disorder (Class 0)	0.50	1.00	0.67	6
Normal (Class 1)	0.00	0.00	0.00	6
accuracy			0.50	12
macro avg	0.25	0.50	0.33	12
weighted avg	0.25	0.50	0.33	12

```


/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to _warn_prf(average, modifier, msg_start, len(result))

cm = confusion_matrix(y_test, predictions)
cm
array([[6, 0],
       [6, 0]])

cm = pd.DataFrame(cm, index = ['0','1'], columns = ['0','1'])

plt.figure(figsize = (10,10))
sns.heatmap(cm,cmap= "Blues", linecolor = 'black', linewidth = 1, annot = True, fmt='',xticklabels = labels,yticklabels = labels)

<Axes: >
```





```
[ ] correct = np.nonzero(predictions == y_test)[0]
    incorrect = np.nonzero(predictions != y_test)[0]
```

Some of the Correctly Predicted Classes



```
[ ] i = 0
    for c in correct[:6]:
        plt.subplot(3,2,i+1)
        plt.xticks([])
        plt.yticks([])
        plt.imshow(x_test[c].reshape(150,150), cmap="gray", interpolation='none')
        plt.title("Predicted Class {},Actual Class {}".format(predictions[c], y_test[c]))
        plt.tight_layout()
        i += 1
```

<ipython-input-22-3b0e8ec19e68>:3: MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated since 3.6 and will be removed two minor releases later; explicitly call ax.remove() as needed.
plt.subplot(3,2,i+1)

Predicted Class 0,Actual Class 0Predicted Class 0,Actual Class 0



Predicted Class 0,Actual Class 0Predicted Class 0,Actual Class 0



Predicted Class 0,Actual Class 0



Some of the Incorrectly Predicted Classes

```
i = 0
for c in incorrect[:6]:
    plt.subplot(3,2,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.imshow(x_test[c].reshape(150,150), cmap="gray", interpolation='none')
    plt.title("Predicted Class {},Actual Class {}".format(predictions[c], y_test[c]))
    plt.tight_layout()
    i += 1
```

<ipython-input-23-d863d2b73988>:3: MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated since 3.6 and will be removed two minor releases later; explicitly call ax.remove() as needed.

Predicted Class 0,Actual Class Predicted Class 0,Actual Class 1



Predicted Class 0,Actual Class Predicted Class 0,Actual Class 1



Predicted Class 0,Actual Class 1

